

ML 710 Project Grading

Project presentations will be judged using the following split:

- (30%) Explanation of ML workload – what model are you parallelizing, on what data, and why is it challenging to parallelize?
- (40%) Explanation of implementation of parallel ML strategy, and how it uses parallel ML aspect instances taught in the course
- (30%) Discussion of experiments and results, and explanation of why your implemented parallel ML strategy did or did not achieve the expected improvement in the metrics of throughput, statistical efficiency and Goodput